

EEL 5737: The Implementation of Distributed File System

Han Wei

Department of Electrical & Computer Engineering
College of Engineering
University of Florida
Gainesville, Florida 32611

Cheng Li

Department of Electrical & Computer Engineering
College of Engineering
University of Florida
Gainesville, Florida 32611

Abstract—In previous assignments, we established a distributed file system. We stored the file data and file metadata into two servers. In this project, we will store the file data into several servers. The new file system can tolerate the server crashes and deal with data corruption. The user can write into the file, even if one file data server is not available.

I. INTRODUCTION

II. DESIGN AND IMPLEMENTATION

A. Basic File Operations

The basic operations of file directory and content are similar to the original functions in previous assignments. However, we need to modify them to meet the new requirements.

B. Redundant Storage

We need to store the data on different servers. We use stateArray to record the reading status of each block. We firstly initialize them as 0. If the data can be read successfully, it is marked 1 and 0, otherwise. The same mechanism goes with the write. The old version data is modified and the status needs to be updated. If the write failed in the servers, the block that is written in will be popped.

```
def distributedWrite(self,
    path,data,offset,fh):
    metafiles = self.getmeta(path)

    startBlk = int(offset/BLOCKSIZE) #
        offset determine the first blockID
        to be read
    startOffset = offset%BLOCKSIZE

    lastBlk =
        int(math.ceil((offset+len(data))/BLOCKSIZE))
    endOffset = (offset+len(data))%BLOCKSIZE

    numblk = lastBlk - startBlk
    start_port = metafiles['startPortID']
    start_port = (start_port +
        startBlk)%self.N
    port = start_port
    state = [0] * (numblk) #initialized as
        0. 1 --- success -1 -----fail
    index_of_block = 0 #this index is to
        record the index of block.
    version = metafiles['version']
    newVersion = version
```

```
for i in range(startBlk, lastBlk):
    if i in range(len(version)):
        newVersion[i] = version[i] + 1
    else:
        newVersion.append(0)
checkpoint = 0 # #counter tp recort how
    much of the data have been written
# The following part is the same logic
    as distributedRead
#1. the old version is modified before
    and need to update
#2. this is the very first new version
for i in range(startBlk,lastBlk):
    if i!=startBlk and i!=lastBlk-1:
        try:
            self.dataPut (port,path,i,data[checkpoint:
                newVersion[i]])
            checkpoint =
                checkpoint+BLOCKSIZE
            state[index_of_block] == 1
        except OSError, e:
            if e.args[0] ==
                errno.ECONNREFUSED:
                state[index_of_block] = -1
                print("NO connection!!!")
                break
    elif i==startBlk:
        try:
            temp = ''
            if newVersion[i]!=0:
                temp =
                    self.dataGet (port,path,i,version[i
                        temp = temp[:startOffset]
                        temp = temp +
                            data[:BLOCKSIZE-startOffset]
                        checkpoint =
                            BLOCKSIZE-startOffset
                        self.dataPut (port,path,i,temp,
                            newVersion[i])
                        state[index_of_block] == 1
        except OSError, e:
            if e.args[0] ==
                errno.ECONNREFUSED:
                state[index_of_block] = -1
                print("NO connection")
                break
    elif i==lastBlk-1: # for the last
        block
        try:
```

```

temp = ''
if newVersion[i]!=0:
    temp =
        self.dataGet(port,path,i,version[i]) as 0. If success, it becomes 1
        temp = temp[endOffset:]
    temp = data[checkpoint:]+temp
    checkpoint = checkpoint +
        len(data[checkpoint:])
    self.dataPut(port,path,i,temp,
        newVersion[i])
    state[index_of_block] == 1
except OSError, e:
    if e.args[0] ==
        errno.ECONNREFUSED:
        state[index_of_block] = -1
        print("No connection..")
        break
#change the server after writing one
    block.
port = (port+1)%self.N
index_of_block += 1

#when the distributedwrite failed in
    some of the servers, pop the block
    just wrote in.
if -1 in state:
    port = metafiles['startPortID']
    for i in range(startBlk,lastBlk):
        if state[i]==1:
            self.dataPop(port,path,i,newVersion[i])
            port = (port+1)%self.N
    return False
else:
    # update the version list
    metafiles['version'] = newVersion
    self.putmeta(path,metafiles)
    return True

def distributedRead(self, path, size,
    offset, fh):
    metafiles = self.getmeta(path)
    content = ''
    # There are two corner cases caused by
        offset and size.
    if offset >= metafiles['st_size']:
        return content
    if (offset+size) > metafiles['st_size']:
        size = metafiles['st_size']-offset

    startBlk = int(offset/BLOCKSIZE) #
        offset determine the first blockID
        to be read
    startOffset = offset%BLOCKSIZE

    # decides which server to read from
        first.
    start_port = metafiles['startPortID']
    start_port = (start_port +
        startBlk)%self.N

    lastBlk =
        int(math.ceil((offset+size)/BLOCKSIZE))
    endOffset = (offset+size)%BLOCKSIZE

    port = start_port # port, assistant
        varieble to traverse the dataservers

```

```

state = [0] * (lastBlk-startBlk)
#need stateArray to record the reading
    status of each block. initialized
    # and it becomes 0 if fails.
#reading each blk from different servers
#since the reading of starting and
    ending block is affected by the
    offset, I operate them separately.
#if it is between, I don't worry about
    offset ,just read them all.
for i in range(startBlk, lastBlk):
    if i != startBlk and i != lastBlk-1:
        content = content +
            self.dataGet(port,path,i,metafiles['vers
    elif i == startBlk:
        content =
            self.dataGet(port,path,i,metafiles['vers
        content = content[startOffset:]
    else:
        rv =
            self.dataGet(port,path,i,metafiles['vers
        content = content +
            rv[:size-len(content))]
    port = (port+1)%self.N

    return content

```

C. Crash Toleration

The file system can recover and resume the serving data from local disk when processes crash and restart. When the data on the server is lost, the file system can recover the data from the replicas on its adjacent servers.

Each server has its own data and the copy of the last server's own data. When the server restarts, the file system will retrieve the history data. If the data on the block is empty, it will get the replica from the next server. If there is still no data on the next server, it will get the replica from other servers.

In case of no history records or hard disk failure, the file system will recover previous data and the copies of the other servers kept by dataserver.

```

def __init__(self,*argv):
    # python dataserver.py 0 3333 4444 5555
        6666
    #The first argument tells which one is the
        current server. e.g., 0 means port
        3333 is the port
    #of current server.
    #In this project, there are 5 arguments,
        the first one is index and the
        following four are ports of the
    #data servers.
    current_server = argv[0]
    N = len(argv)-1
    Ports_dataServ = argv[1:]
    port = str(Ports_dataServ[current_server])
    print ('dataserver',port,'is online...')
    # Retrieve history data
    # self.data is a dict of dicts that has
        REDUNDANCY ports:data pairs
    # Actually, each server has its own data
        and the copy of last server's own data

```

```

# {}, the filename is dataserver_port

self.data =
    shelve.open('dataserver_'+port,
        writeback=True)

# connect to the nextServer for recovery
nextPort =
    str(Ports_dataServ[(current_server+1)%N])
self.nextServer =
    xmlrpclib.ServerProxy('http://localhost:'
        + nextPort)

self.ports = {}
self.servers = {}
for i in range(REDUNDANCY):
    self.ports[i] =
        str(Ports_dataServ[(current_server-i)%N])
        # ports[0] is the port of current
        # dataserver
    if i!= 0:
        self.servers[self.ports[i]] =
            xmlrpclib.ServerProxy('http://localhost:'
                + self.ports[i])
#connect to the all the other dataservers
#check whether the data dic is empty. If
#it is empty, initialize them. get the
#replica from nextserver.
#Otherwise, get the replica from other
#servers.
if self.data!= {}:
    print("Loaded previous data from server
        backup file")
else:
    for i in range(REDUNDANCY):
        self.data[self.ports[i]] = {}
    # 2 cases:
    # 1)No history records:
    # 2)Hard disk fails
    # Anyway, recover previous data and the
    # copies of other servers kept by this
    # dataserver
    for i in self.ports:
        # Communicate with the next server to
        # retrieve its own data
        # 1)recover self.data[port]: talk with
        # next server
        p = self.ports[i]
        if p == port:
            # Assume only one server crash at a
            # time.
            # No need to retry if failed to
            # connect with server
            try:
                replica =
                    self.nextServer.getReplica(Binary(port))
                self.data[p] =
                    pickle.loads(replica.data)
                print ('Successfully loaded previous
                    data from dataserver:', nextPort)
            except socket.error:
                print ('Recovering previous data is
                    passed')
            pass # Based on the assumption, this
            # only happens when starting to
            # run all dataservers at the
            # beginning. In this case, skip

```

```

        getReplica by ignore the socket
        error
    else:
        # 2)recover other copies from the
        # original dataservers
        try:
            replica =
                self.servers[p].getReplica(Binary(p))
            self.data[p] =
                pickle.loads(replica.data)
            print ("Successfully loaded previous
                copy from dataserver:", p)
        except socket.error:
            print ("Recovering previous copy is
                passed")
            pass

```

D. Corruption Protection

For test, a corrupt function is made. The function will corrupt at least one byte of any blocks of the latest file. A block and some bytes are chosen randomly. The chosen bytes will be messed intentionally.

```

def corrupt(self,path): # corrupt the latest
    data
    # This function will corrupt at least one
    # byte of any blocks of the file
    rv = False
    p = []
    for i in self.ports:
        if path in self.data[self.ports[i]]:
            p.append(i)
    if len(p)>0:
        s = np.random.randint(0,len(p))
        s = p[s]

        # randomly pick up a block
        IDlist =
            self.data[self.ports[s]][path].keys()
        ID = np.random.randint(0,len(IDlist))
        ID = IDlist[ID]
        # randomly pick up a byte
        version =
            len(self.data[self.ports[s]][path][ID])-1
        mes =
            self.data[self.ports[s]][path][ID][version]
        byte = np.random.randint(0,len(mes)-1)
        i = 1
        while(mes[byte]==str(i)):
            i = i+1
        mes = mes[:byte] + str(i) + mes[byte+1:]
        self.data[self.ports[s]][path][ID][version]
            = mes
        rv = True
        print('Information:')
        print ('Corrupted file:', path)
        print ('Corrupted blks: ', ID)
    else:
        print("Path is unfound.")
    return rv

```

For corruption protection, the checksum is needed. We use cyclic redundancy check (CRC), which is widely used in all

```
def checksum(self,data):
    crc = binascii.crc32(data) & 0xffffffff
    return "{0:x}".format(crc).rjust(8,'0') #
        take it as a 8-byte string
```

```

crc1 = message[len(message)-8:]
crc = self.checksum(data[i])
if (crc==crc1):
    server_states[i+REDUNDANCY] = 1

```

The status array can record the status of write. If the write fails because of server failure or so, the block will pop.

```
#when the distributedwrite failed in some of
    the servers, pop the block just wrote in.
if -1 in state:
    port = metafiles['startPortID']
    for i in range(startBlk,lastBlk):
        if state[i]==1:
            self.dataPop(port,path,i,newVersion[i])
            port = (port+1)%self.N
        return False
else:
    # update the version list
    metafiles['version'] = newVersion
    self.putmeta(path,metafiles)
    return True
```

A. Creating Directories

```
cheng@cheng-Allienware-17-R3:~/Desktop/pocsd project/3/fusepy/fusenount$ ls
cheng@cheng-Allienware-17-R3:~/Desktop/pocsd project/3/fusepy/fusenount$ mkdir 1
cheng@cheng-Allienware-17-R3:~/Desktop/pocsd project/3/fusepy/fusenount$ mkdir 2
cheng@cheng-Allienware-17-R3:~/Desktop/pocsd project/3/fusepy/fusenount$ mkdir 3
cheng@cheng-Allienware-17-R3:~/Desktop/pocsd project/3/fusepy/fusenount$ ls
```

B. Creating Files

```
chengcheng-Allenware-17-R3:~/Desktop/pocsd project/3/fusepy/fusenount/1$ ls
test test2
chengcheng-Allenware-17-R3:~/Desktop/pocsd project/3/fusepy/fusenount/1$ cd ..
chengcheng-Allenware-17-R3:~/Desktop/pocsd project/3/fusepy/fusenount$ ls
2
chengcheng-Allenware-17-R3:~/Desktop/pocsd project/3/fusepy/fusenount$ mv 1 2
chengcheng-Allenware-17-R3:~/Desktop/pocsd project/3/fusepy/fusenount$ cd 2
chengcheng-Allenware-17-R3:~/Desktop/pocsd project/3/fusepy/fusenount/2$ ls
chengcheng-Allenware-17-R3:~/Desktop/pocsd project/3/fusepy/fusenount/2$ cd 1
chengcheng-Allenware-17-R3:~/Desktop/pocsd project/3/fusepy/fusenount/2/1$ ls
test test2
chengcheng-Allenware-17-R3:~/Desktop/pocsd project/3/fusepy/fusenount/2/1$ cat test
111111111111111111111111
chengcheng-Allenware-17-R3:~/Desktop/pocsd project/3/fusepy/fusenount/2/1$ ls
```

C. Creating Links

```
chengcheng-Allenware-17-R3:~/Desktop/pcosd project/3/fusepy/fusenount$ ls -l
lrwxrwxrwx 1 root root 0 Dec 3 21:13:33 linked
chengcheng-Allenware-17-R3:~/Desktop/pcosd project/3/fusepy/fusenount$ ln -s linked mysymlink
chengcheng-Allenware-17-R3:~/Desktop/pcosd project/3/fusepy/fusenount$ ln -l linked mysymlink
ln: invalid option -- 'l'
Try 'ln --help' for more information.
chengcheng-Allenware-17-R3:~/Desktop/pcosd project/3/fusepy/fusenount$ ls -l
lrwxrwxrwx 1 root root 0 Dec 3 21:13:33 linked
lrwxrwxrwx 1 root root 0 Dec 3 19:09 mysymlink -> linked
chengcheng-Allenware-17-R3:~/Desktop/pcosd project/3/fusepy/fusenount$
```

Fig. 3: Creating links.

Making server 4444 crash manually, and recovering the server later.

```

py 1 3333 4444 5555 6666
dataserver 4444 is online...
Loaded previous data from server backup file
127.0.0.1 - - [03/Dec/2017 21:47:59] "POST /RPC2 HTTP/1.1" 200 -
127.0.0.1 - - [03/Dec/2017 21:47:59] "POST /RPC2 HTTP/1.1" 200 -
127.0.0.1 - - [03/Dec/2017 21:47:59] "POST /RPC2 HTTP/1.1" 200 -
127.0.0.1 - - [03/Dec/2017 21:47:59] "POST /RPC2 HTTP/1.1" 200 -
127.0.0.1 - - [03/Dec/2017 22:07:25] "POST /RPC2 HTTP/1.1" 200 -
127.0.0.1 - - [03/Dec/2017 22:07:25] "POST /RPC2 HTTP/1.1" 200 -
127.0.0.1 - - [03/Dec/2017 22:07:25] "POST /RPC2 HTTP/1.1" 200 -
127.0.0.1 - - [03/Dec/2017 22:07:27] "POST /RPC2 HTTP/1.1" 200 -
127.0.0.1 - - [03/Dec/2017 22:07:27] "POST /RPC2 HTTP/1.1" 200 -
127.0.0.1 - - [03/Dec/2017 22:07:34] "POST /RPC2 HTTP/1.1" 200 -
127.0.0.1 - - [03/Dec/2017 22:07:34] "POST /RPC2 HTTP/1.1" 200 -
127.0.0.1 - - [03/Dec/2017 22:07:38] "POST /RPC2 HTTP/1.1" 200 -
127.0.0.1 - - [03/Dec/2017 22:07:38] "POST /RPC2 HTTP/1.1" 200 -
127.0.0.1 - - [03/Dec/2017 22:07:51] "POST /RPC2 HTTP/1.1" 200 -
127.0.0.1 - - [03/Dec/2017 22:07:51] "POST /RPC2 HTTP/1.1" 200 -
127.0.0.1 - - [03/Dec/2017 22:07:51] "POST /RPC2 HTTP/1.1" 200 -
127.0.0.1 - - [03/Dec/2017 22:07:51] "POST /RPC2 HTTP/1.1" 200 -
127.0.0.1 - - [03/Dec/2017 22:07:51] "POST /RPC2 HTTP/1.1" 200 -
127.0.0.1 - - [03/Dec/2017 22:08:02] "POST /RPC2 HTTP/1.1" 200 -
127.0.0.1 - - [03/Dec/2017 22:08:02] "POST /RPC2 HTTP/1.1" 200 -
C:\server>chaged manually...
cheng@cheng-allenware-17-R3:~/Desktop/pocsd project/3/fusepy$

```

Fig. 4: The server crashes manually.

```
cheng@cheng-Allenware-17-R3:~/Desktop/pocsd project/3/fusepy/fusemount/1$ ls  
crashtest test test2  
cheng@cheng-Allenware-17-R3:~/Desktop/pocsd project/3/fusepy/fusemount/1$ echo  
!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!! > crashtest  
bash: echo: write error: Bad address  
cheng@cheng-Allenware-17-R3:~/Desktop/pocsd project/3/fusepy/fusemount/1$
```

Fig. 5: Write into the crashed server.

[illegible]

Fig. 6: Recover the server.

[illegible]

Fig. 7: Recover the server.

V. CONCLUSION

VI. CONTRIBUTION

VII. APPENDIX

```
distributedFS.py
dataserver.py
metaserver.py
```

REFERENCES

- [1] On-line crc calculation and free library. [Online]. Available: <https://www.lammertbies.nl/comm/info/crc-calculation.html>